



hcHyper

陈岳 清华大学计算机科学与技术系



Overview

hcHyper: <https://github.com/cylindrical2002/hcHyper>

hcHyper 项目实现了基于 ArceOS 提供的硬件支持下的 x86(Intel) , ARM , RISC-V 三大架构的 CPU 和 Memory 的虚拟化支持。CPU 支持部分实现在 `hypervisor/crates/hypercraft` 中, Memory 部分实现在 `hypervisor/crates` 下的 `guest_page_table` 和 `guest_page_table_entry` 中。

目前能够支持在含 KVM 的 Intel PC 中运行 Hypervisor 虚拟化 NimbOS 操作系统, 同时支持在 QEMU 中运行 ARM 和 RISC-V 两种架构下的 Hypervisor , 分别能够支持 NimbOS 操作系统和 Linux 操作系统。



Predessor

hcHyper 仓库的三大前身分别是 **RVM-Tutorial** , **rHyper** , **hypercraft**

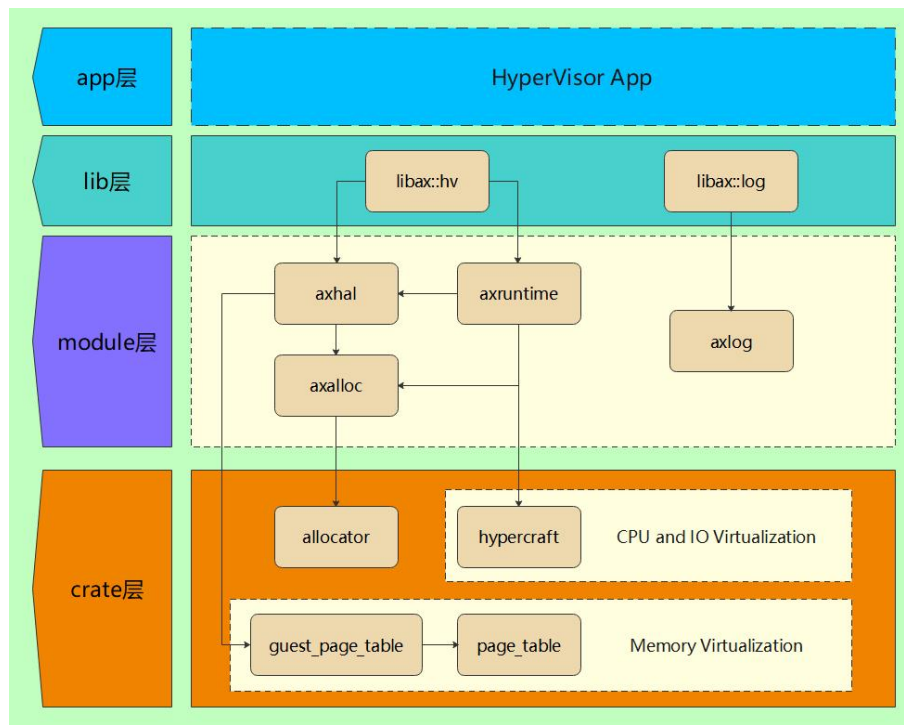
RVM-Tutorial 是一个用 **RUST** 语言编写的 **Type-1 Hypervisor** , 支持在 **x86(Intel)** 的 **PC** 上运行, 并支持 **x86(Intel)** 下的 **CPU** 和 **Memory** 的虚拟化, 能够支持 **NimbOS** 操作系统。

rHyper 是一个用 **RUST** 语言编写的 **Type-1 Hypervisor** , 支持在 **QEMU** 上运行, 并支持 **ARM-v8** 下的 **CPU** 、 **Memory** 、 **IO** 的虚拟化, 能够支持 **NimbOS \ Linux \ rCore \ ArceOS** 操作系统, 并且能够支持多核虚拟化。

hypercraft 是一个用 **RUST** 语言编写的 **Type-2 Hypervisor** , 支持在 **QEMU** 上运行, 并支持 **RISC-V** 下的 **CPU** 、 **Memory** 、 **IO** 的虚拟化, 能够支持 **NimbOS** 操作系统。是首个结合 **ArceOS** 框架的操作系统。



Design





Result

```
Mapping physical memory: [0xfffff800068f000, 0xfffff8001000000)
Mapping MMIO: [0xfffff80fec00000, 0xfffff80fec01000)
Mapping MMIO: [0xfffff80fed00000, 0xfffff80fed01000)
Mapping MMIO: [0xfffff80fee00000, 0xfffff80fee01000)
Initializing drivers...
Initializing local APIC...
Initializing HPET...
HPET: 100.000000 MHz, 64-bit, 3 timers
Calibrated TSC frequency: 4007.558 MHz
Calibrated LAPIC frequency: 1000.119 MHz
Initializing task manager...
/**** APPS ****
cyclictest
exit
fantastic_text
forktest
forktest2
forktest_simple
forktest_simple_c
forktree
hello_c
hello_world
matrix
sleep
sleep_simple
stack_overflow
thread_simple
user_shell
usertests
yield
*****
Running tasks...
test kernel task: pid = TaskId(2), arg = 0xdead
test kernel task: pid = TaskId(3), arg = 0xbeef
Rust user shell
>>

Initializing frame allocator at: [PA:0x40518000, PA:0x48000000)
Mapping .text: [0xffff000040080000, 0xffff000040094000)
Mapping .rodata: [0xffff000040094000, 0xffff00004009b000)
Mapping .data: [0xffff00004009b000, 0xffff000040110000)
Mapping .bss: [0xffff000040115000, 0xffff000040518000)
Mapping boot stack: [0xffff00004011000, 0xffff000040115000)
Mapping physical memory: [0xffff000040518000, 0xffff000048000000)
Mapping MMIO: [0xffff000090000000, 0xffff000090010000)
Mapping MMIO: [0xffff000080000000, 0xffff000080200000)
Initializing drivers...
Initializing task manager...
/**** APPS ****
cyclictest
exit
fantastic_text
forktest
forktest2
forktest_simple
forktest_simple_c
forktree
hello_c
hello_world
matrix
sleep
sleep_simple
stack_overflow
thread_simple
user_shell
usertests
yield
*****
Running tasks...
test kernel task: pid = TaskId(2), arg = 0xdead
test kernel task: pid = TaskId(3), arg = 0xbeef
Rust user shell
>>

~ # ls
[ 11.380783 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.265126] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.382091 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.266422] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.385147 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.269498] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.386424 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.270795] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.387668 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.271779] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.388536 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.272886] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.389758 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.273916] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.390817 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.275791] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.392703 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.277284] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.394222 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.278948] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.395884 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.280280] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.397217 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.281582] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.398505 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.283309] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.400250 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.284671] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.401703 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.286113] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])
[ 11.403038 0 hypercraft::arch:vm:83] Remote fence is not supported
[ 11.287346] __sbi_rfence_v02_call: hbase = [0] hmask = [0x1] failed (error [-524])

bin      etc      lost+found /sbin      test
dev      linuxrc  proc       sys       usr
~ #
```



CPU Virtualization Basic

我们可以先温习一下有关于 **CPU** 虚拟化技术的知识，这些知识对于三大架构，乃至所有的不模拟硬件的 **Hypervisor** 软件都适用。

hypervisor init: 初始化一个 **hypervisor**，包括初始化运行 **hypervisor** 的 **CPU** 基本启动信息，以及 **hypervisor** 本身的设置。

CPU allocation: 分配 **Guest** 执行的 **VCPU**，并分配对应的 **PerCPU**

VCPU: 由 **hypervisor** 虚拟出来的，运行 **guest** 软件所需的每 **CPU** 的私有状态。类似于传统 **OS** 中的线程，一个 **guest OS** 可具有多个 **vCPU**，**vCPU** 数量与物理 **CPU** 数量无关。

PerCPU: **hypervisor** 管理的所有物理 **CPU**。

Memory allocation: 分配 **Guest** 执行时的所有 **VCPU** 中的内存空间以及对应的真实物理内存。并建立 **guest_page_table** 保存对应的页表指针等等。

VM entry: 从 **host** 模式切换到 **guest** 模式，开始执行 **guest** 软件的代码。

VM exit: 从 **guest** 模式切换回 **host** 模式，开始执行 **hypervisor** 的代码。

Device virtualization: 虚拟化 **guest** 软件执行的设备，包含直通，模拟等几种方式。



CPU Virtualization Basic

```
/// Trait for VM struct.
pub trait VmTrait<H: HyperCraftHal, G: GuestMemoryInterface> {
    /// Create a new VM with `vcpus` vCPUs and `gpt` as the guest page table.
    fn new(vcpus: VmCpus<H>, gpt: G) -> HyperResult<Self>
    where
        Self: Sized;

    /// Initialize `VCpu` by `vcpu_id`.
    fn init_vcpu(&mut self, vcpu_id: usize);

    /// Run the host VM's vCPU with ID `vcpu_id`. Does not return.
    fn run(&mut self, vcpu_id: usize);
}
```

```
pub trait RvmHal: Sized {
    /// Allocates a 4K-sized contiguous physical page, returns its physical address.
    fn alloc_page() -> Option<HostPhysAddr>;
    /// Deallocates the given physical page.
    fn dealloc_page(paddr: HostPhysAddr);
    /// Converts a physical address to a virtual address which can access.
    fn phys_to_virt(paddr: HostPhysAddr) -> HostVirtAddr;
    /// Converts a virtual address to the corresponding physical address.
    fn virt_to_phys(vaddr: HostVirtAddr) -> HostPhysAddr;
    /// VM-Exit handler.
    fn vmexit_handler(vcpu: &mut crate::RvmVcpu<Self>);
}
```



PerCPU

```
/// Trait for PerCpu struct.
pub trait PerCpuTrait<H: HyperCraftHal> {
    /// Initializes the `PerCpu` structures for each CPU. This (the boot CPU's) per-CPU
    /// area is initialized and loaded into TP as well.
    fn init(boot_hart_id: usize, stack_size: usize) -> HyperResult<()>;

    /// Initializes the `thread` pointer to point to PerCpu data.
    fn setup_this_cpu(hart_id: usize) -> HyperResult<()>;

    /// Create a `VCpu`, set the entry point to `entry` and bind this vcpu into the current CPU.
    fn create_vcpu(&mut self, vcpu_id: usize, entry: GuestPhysAddr) -> HyperResult<VCpu<H>>;

    /// Returns this CPU's `PerCpu` structure.
    fn this_cpu() -> &'static mut Self;
}
```




VCPU

```
/// Trait for Vcpu struct.
pub trait VCpuTrait {
    /// Create a new vCPU
    fn new(vcpu_id: usize, entry: GuestPhysAddr) -> Self;

    /// Runs this vCPU until traps.
    fn run(&mut self) -> VmExitInfo;

    /// Gets one of the vCPU's general purpose registers.
    fn get_gpr(&self, index: GprIndex);

    /// Set one of the vCPU's general purpose register.
    fn set_gpr(&mut self, index: GprIndex, val: usize);

    /// Gets the vCPU's id.
    fn vcpu_id(&self) -> usize;
}
```



X86 Init

1. x86 的基本启动设置

1. 检查硬件的虚拟化支持

1. 初始化 **VMXON** 空间

1. 执行 **VMXON**

1. 初始化 **VMCS**



X86 Init

```
clear_bss();
arch::init_early();
println!("{}", LOGO);
println!["\n\
arch = {}\n\
build_mode = {}\n\
log_level = {}\n\
",
option_env!("ARCH").unwrap_or(""),
option_env!("MODE").unwrap_or(""),
option_env!("LOG").unwrap_or(""),
];

mm::init_heap_early();
logging::init();
info!("Logging is enabled.");

arch::init();
mm::init();
INIT_OK.store(val: true, order: Ordering::SeqCst);
println!("Initialization completed.\n");
```

```
unsafe {
    // Enable VMX using the VMXE bit.
    Cr4::write(flags: Cr4::read() | Cr4Flags::VIRTUAL_MACHINE_EXTENSIONS);
    // Execute VMXON.
    vmx::vmxon(self.vmx_region.phys_addr() as _)?;
}
info!("[RVM] succeeded to turn on VMX.");
```

```
pub fn has_hardware_support() -> bool {
    if let Some(feature: FeatureInfo) = CpuId::new().get_feature_info() {
        feature.has_vmx()
    } else {
        false
    }
}
```

```
// Get VMCS revision identifier in IA32_VMX_BASIC MSR.
let vmx_basic: VmxBasic = VmxBasic::read();
if vmx_basic.region_size as usize != crate::mm::PAGE_SIZE {
    return rvm_err!(Unsupported);
}
if vmx_basic.mem_type != VmxBasic::VMX_MEMORY_TYPE_WRITE_BACK {
    return rvm_err!(Unsupported);
}
if vmx_basic.is_32bit_address {
    return rvm_err!(Unsupported);
}
if !vmx_basic.io_exit_info {
    return rvm_err!(Unsupported);
}
if !vmx_basic.vmx_flex_controls {
    return rvm_err!(Unsupported);
}
self.vmxcs_revision_id = vmx_basic.revision_id;
self.vmx_region = VmxRegion::new(vmx_basic.revision_id, shadow_indicator: false)?;
```

```
fn setup_vmcs(&mut self, entry: GuestPhysAddr, ept_root: HostPhysAddr) -> RvmResult {
    let paddr: u64 = self.vmxcs.phys_addr() as u64;
    unsafe {
        vmx::vmclear(addr: paddr)?;
        vmx::vmptwld(addr: paddr)?;
    }
    self.setup_vmcs_host()?;
    self.setup_vmcs_guest(entry)?;
    self.setup_vmcs_control(ept_root)?;
    Ok(())
}
```



ARM Init

1. ARM 的基本启动设置

1. 检查硬件的虚拟化支持

1. 设置 HCR_EL2 寄存器



ARM Init

```
arch::init();
device::init();
info!("Hello World from cpu {}", cpu_id);
// Safety: Modify to usize is atomic; there is most one writer at the same time.
unsafe {
    while GUEST_ENTRIES[cpu_id] == 0 {
        trace!("secondary cpu {} waiting", cpu_id);
    }
    info!(
        "secondary cpu {} will run a vcpu with entry 0x{:x}",
        cpu_id, GUEST_ENTRIES[cpu_id]
    );
    hv::run(cpu_id, entry: GUEST_ENTRIES[cpu_id], psci_context: PSCI_CONTEXT[cpu_id]);
}
```

```
HCR_EL2.write(
    field: HCR_EL2::VM::Enable + HCR_EL2::RW::EL1IsAarch64 + HCR_EL2::AMO::SET + HCR_EL2::FMO::SET
);
Ok(())
```



RISC-V Init

1. RISC-V 的基本启动设置

1. 检查硬件的虚拟化支持

1. 初始化 `hypervisor` 寄存器



RISC-V Init

```
// Detect if hypervisor extension exists on current hart environment
//
// This function tries to read hgatp and returns false if the read operation failed.
pub fn detect_h_extension() -> bool {
    // run detection by trap on csrr instruction.
    let ans = with_detect_trap(0, || unsafe {
        asm!("csrr {}, 0x680", out(reg) _, options(nomem, nostack)); // 0x680 => hgatp
    });
    // return the answer from output flag. 0 => success, 2 => failed, illegal instruction
    ans != 2
}
```

```
let mut regs = riscv::register::read();
// Set hstatus
let mut hstatus = LocalRegisterCopy::<usize, hstatus::Register>::new(
    riscv::register::hstatus::read().bits(),
);
hstatus.modify(hstatus::spv::Supervisor);
// Set SPVP bit in order to accessing VS-mode memory from HS-mode.
hstatus.modify(hstatus::spvp::Supervisor);
CSR.hstatus.write_value(hstatus.get());
regs.guest_regs.hstatus = hstatus.get();

// Set sstatus
let mut sstatus = sstatus::read();
sstatus.set_spp(sstatus::SPP::Supervisor);
regs.guest_regs.sstatus = sstatus.bits();

regs.guest_regs.gprs.set_reg(GprIndex::A0, 0);
regs.guest_regs.gprs.set_reg(GprIndex::A1, 0x9000_0000);
```



x86 launch

```
asm!(  
    "mov    [rdi + {host_stack_top}], rsp", // save current RSP to Vcpu::host_stack_top  
    "mov    rsp, rdi",                      // set RSP to guest regs area  
    restore_regs_from_stack!(),  
    "vmlaunch",  
    "jmp    {failed}",  
    host_stack_top = const size_of::<GeneralRegisters>(),  
    failed = sym Self::vmx_entry_failed,  
    options(noreturn),  
)
```




arm launch

```
asm!(  
    "mov    x28, sp",  
    "str    x28, [x0, {host_stack_top}]", // save current SP to Vcpu::host_stack_top  
    "mov    sp, x0", // set SP to guest regs area  
    restore_regs_from_stack!(),  
    "eret",  
    "bl     {failed}",  
    host_stack_top = const size_of::<GeneralRegisters>() + 3 * size_of::<u64>(),  
    failed = sym Self::vmentry_failed,  
    options(noreturn),  
)
```



risc-v launch

```
/* Set start point */  
ld    t1, ({guest_sepc})(a0)  
csrw  sepc, t1  
  
/* Set stvec so that hypervisor resumes after the sret when the guest exits. */  
la    t1, _guest_exit  
csrrw t1, stvec, t1  
sd    t1, ({hyp_stvec})(a0)  
  
/* Save sscratch and replace with pointer to GuestInfo. */  
csrrw t1, sscratch, a0  
sd    t1, ({hyp_sscratch})(a0)
```

```
sret
```



x86 exit-handler

1. `CPUID` 指令的处理

1. `VMCALL` 对 `hypervisor` 的调用

1. `EPT VIOLATION`



x86 exit-handler

```
let exit_info: VmxExitInfo = vcpu.exit_info()?;
trace!("VM exit: {:#x?}", exit_info);

if exit_info.entry_failure {
    panic!("VM entry failed: {:#x?}", exit_info);
}

let res: Result<(), RvmError> = match exit_info.exit_reason {
    VmxExitReason::CPUID => handle_cpuid(vcpu),
    VmxExitReason::VMCALL => handle_hypercall(vcpu),
    VmxExitReason::EPT_VIOLATION => handle_ept_violation(vcpu, exit_info.guest_rip),
    _ => panic!(
        "Unhandled VM-Exit reason {:?}: \n{:#x?}",
        exit_info.exit_reason, vcpu
    ),
};

if res.is_err() {
    panic!(
        "Failed to handle VM-exit {:?}: \n{:#x?}",
        exit_info.exit_reason, vcpu
    );
}
```

Ok({})

Yuekai Jia, 7 months ago • Step 3: run a simple guest



arm exit-handler

1. hypercall HVC64 指令

1. DataAbort \ InstAbort



arm exit-handler

```
let exit_info: ArmExitInfo = vcpu.exit_info()?;
// if vcpu.cpu_id != 0 {
// println!("cpu {} exit", vcpu.cpu_id);
// }
// debug!("VM exit: {:#x?}", exit_info);

let res: Result<(), RvmError> = match exit_info.exit_reason {
    Some(ESR_EL2::EC::Value::HVC64) => handle_hyercall(vcpu),
    Some(ESR_EL2::EC::Value::InstrAbortLowerEL)
    | Some(ESR_EL2::EC::Value::InstrAbortCurrentEL) => handle_iabt(vcpu),
    Some(ESR_EL2::EC::Value::DataAbortLowerEL)
    | Some(ESR_EL2::EC::Value::DataAbortCurrentEL) => handle_dabt(vcpu),
    _ => panic!(
        "Unhandled VM-Exit reason {:?}:\n{:#x?}",
        exit_info.exit_reason.unwrap() as u64,
        vcpu
    ),
};

if res.is_err() {
    panic!(
        "Failed to handle VM-exit {:?}:\n{:#x?}",
        exit_info.exit_reason.unwrap() as u64,
        vcpu
    );
}
```



risc-v exit-handler

1. hypercall ECALL 指令

1. PAGEFAULT

1. TIMER GuestOS 时钟中断



risc-v exit-handler

```
let scause = scause::read();
use scause::{Exception, Interrupt, Trap};
// Use SCAUSE to find a VmExitInfo for VM to deal with
match scause.cause() {
  Trap::Exception(Exception::VirtualSupervisorEnvCall) => {
    let sbi_msg = SbiMessage::from_regs(regs.guest_regs.gprs.a_regs()).ok();
    VmExitInfo::Ecall(sbi_msg)
  }
  Trap::Interrupt(Interrupt::SupervisorTimer) => VmExitInfo::TimerInterruptEmulation,
  Trap::Interrupt(Interrupt::SupervisorExternal) => {
    VmExitInfo::ExternalInterruptEmulation
  }
  Trap::Exception(Exception::LoadGuestPageFault)
  | Trap::Exception(Exception::StoreGuestPageFault) => {
    let fault_addr = regs.trap_csrs.htval << 2 | regs.trap_csrs.stval & 0x3;
    // debug!(
    //   "fault_addr: {:#x}, htval: {:#x}, stval: {:#x}, sepc: {:#x}, scause: {:?}",
    //   fault_addr,
    //   regs.trap_csrs.htval,
    //   regs.trap_csrs.stval,
    //   regs.guest_regs.sepc,
    //   scause.cause()
    // );
    VmExitInfo::PageFault {
      fault_addr,
      // Note that this address is not necessarily guest virtual as the guest may or
      // may not have 1st-stage translation enabled in VSATP. We still use GuestVirtAddr
      // here though to distinguish it from addresses (e.g. in HTVAL, or passed via a
      // TEECALL) which are exclusively guest-physical. Furthermore we only access guest
      // instructions via the HLVX instruction, which will take the VSATP translation
      // mode into account.
      fault_pc: regs.guest_regs.sepc,
      inst: regs.trap_csrs.htinst as u32,
      priv_level: PrivilegeLevel::from_hstatus(regs.guest_regs.hstatus),
    }
  }
}
=> {
  panic!(
    "Unhandled trap: {:?}, sepc: {:#x}, stval: {:#x}",
    scause.cause(),
    regs.guest_regs.sepc,
    regs.trap_csrs.stval
  );
}
```




Memory Virtualization Basic

我们同样先温习一下有关于 **Memory** 虚拟化技术的知识，这些知识对于三大架构，乃至所有的不模拟硬件的 **Hypervisor** 软件都适用。

Memory allocation: 分配 **Guest** 执行时的所有 **VCPU** 中的内存空间以及对应的真实物理内存。并建立 **guest_page_table** 保存对应的页。表指针等等

guest page table: 支持 **guest physaddr** 到真实的物理地址的地址转换的页表

guest physaddr: 用户认为自己所在的地址

two stage translation: 二级页表转换，即 **guest app** 的 **guest virtaddr** 向 **guest os** 的 **guest physaddr** 再向真实的物理地址的转换。



ArceOS PageTable

```
const ENTRY_COUNT: usize = 512;

const fn p4_index(vaddr: VirtAddr) -> usize {
    (vaddr.as_usize() >> (12 + 27)) & (ENTRY_COUNT - 1)
}

const fn p3_index(vaddr: VirtAddr) -> usize {
    (vaddr.as_usize() >> (12 + 18)) & (ENTRY_COUNT - 1)
}

const fn p2_index(vaddr: VirtAddr) -> usize {
    (vaddr.as_usize() >> (12 + 9)) & (ENTRY_COUNT - 1)
}

const fn p1_index(vaddr: VirtAddr) -> usize {
    (vaddr.as_usize() >> 12) & (ENTRY_COUNT - 1)
}
```



ArceOS PageTable

```
impl<M: PagingMetaData, PTE: GenericPTE, IF: PagingIf> PageTable64<M, PTE, IF> {  
    /// For Construct it out of the crate  
    pub fn new(_root_paddr: PhysAddr, _intrm_tables: Vec<PhysAddr>) -> Self { ...  
  
        /// Creates a new page table instance or returns the error.  
        ///  
        /// It will allocate a new page for the root page table.  
    pub fn try_new() -> PagingResult<Self> { ...  
  
        /// Returns the physical address of the root page table.  
    pub const fn root_paddr(&self) -> PhysAddr { ...  
  
        /// Maps a virtual page to a physical frame with the given `page_size` ...  
    pub fn map(...  
    ) -> PagingResult { ...  
  
        /// Unmaps the mapping starts with `vaddr`. ...  
    pub fn unmap(&mut self, vaddr: VirtAddr) -> PagingResult<(PhysAddr, PageSize)> { ...  
  
        /// Query the result of the mapping starts with `vaddr`. ...  
    pub fn query(&self, vaddr: VirtAddr) -> PagingResult<(PhysAddr, MappingFlags, PageSize)> { ...  
  
        /// Map a contiguous virtual memory region to a contiguous physical memory ...  
    pub fn map_region(...  
    ) -> PagingResult { ...  
  
        /// Unmap a contiguous virtual memory region. ...  
    pub fn unmap_region(&mut self, vaddr: VirtAddr, size: usize) -> PagingResult { ...  
  
        /// Walk the page table recursively. ...  
    pub fn walk<F>(&self, limit: usize, func: &F) -> PagingResult  
    where  
        F: Fn(usize, usize, VirtAddr, &PTE),  
    { ...  
    }  
}
```

ArceOS PageTable

```
/// The low-level **OS-dependent** helpers that must be provided for
/// [PageTable64].
pub trait PagingIf: Sized {
    /// Request to allocate a 4K-sized physical frame.
    fn alloc_frame() -> Option<PhysAddr>;

    /// Request to free a allocated physical frame.
    fn dealloc_frame(paddr: PhysAddr);

    /// Returns a virtual address that maps to the given physical address.
    ///
    /// Used to access the physical memory directly in page table implementation.
    fn phys_to_virt(paddr: PhysAddr) -> VirtAddr;
}
```



Guest PageTable

```
fn alloc_guest_page_table() -> PagingResult<PhysAddr> {  
    if let Some(paddr) = IF::alloc_frames(4) {  
        let ptr = IF::phys_to_virt(paddr).as_mut_ptr();  
        unsafe { core::ptr::write_bytes(ptr, 0, PAGE_SIZE_4K * 4) };  
        Ok(paddr)  
    } else {  
        Err(PagingError::NoMemory)  
    }  
}
```



Guest PageTable

You, 2 weeks ago | 1 author (You)

```
pub trait GuestPagingIf: PagingIf {  
    /// Request to allocate `page_nums` 4K-sized physical frame.  
    #[cfg(target_arch = "riscv64")]  
    fn alloc_frames(page_nums: usize) -> Option<PhysAddr>;  
  
    /// Request to free `page_nums` 4K-sized physical frame.  
    #[cfg(target_arch = "riscv64")]  
    fn dealloc_frames(paddr: PhysAddr, page_nums: usize);  
}
```